Congratulations! You passed!

TO PASS 80% or higher

Keep Learning

GRADE
93.33%

Bird recognition in the city of Peacetopia (case study)

LATEST SUBMISSION GRADE

93.33%

1. ## Problem Statement

1 / 1 point

This example is adapted from a real production application, but with details disguised to protect confidentiality.

You are a famous researcher in the City of Peacetopia. The people of Peacetopia have a common characteristic: they are afraid of birds. To save them, you have **to build an algorithm that will detect any bird flying over Peacetopia** and alert the population.

The City Council gives you a dataset of 10,000,000 images of the sky above Peacetopia, taken from the city's security cameras. They are labelled:

- y = 0: There is no bird on the image
- y = 1: There is a bird on the image

Your goal is to build an algorithm able to classify new images taken by security cameras from Peacetopia.

There are a lot of decisions to make:

- What is the evaluation metric?
- How do you structure your data into train/dev/test sets?

Metric of success

The City Council tells you that they want an algorithm that

1. Has high accuracy
2. Runs quickly and takes only a short time to classify a new image.
3. Can fit in a small amount of memory, so that it can run in a small processor that the city will attach to many different security cameras.

Note: Having three evaluation metrics makes it harder for you to quickly choose between two different algorithms, and will slow down the speed with which your team can iterate. True/False?

- (●) True
- ( ) False

✓ **Correct**

2. After further discussions, the city narrows down its criteria to: **1 / 1 point**

- "We need an algorithm that can let us know a bird is flying over Peacetopia as accurately as possible."
- "We want the trained model to take no more than 10sec to classify a new image."
- "We want the model to fit in 10MB of memory."

If you had the three following models, which one would you choose?

Test Accuracy	Runtime	Memory size
97%	1 sec	3MB

Test Accuracy	Runtime	Memory size
99%	13 sec	9MB

Test Accuracy	Runtime	Memory size
97%	3 sec	2MB

Test Accuracy	Runtime	Memory size
98%	9 sec	9MB

Correct

Correct! As soon as the runtime is less than 10 seconds you're

2. After further discussions, the city narrows down its criteria to: **1 / 1 point**

- "We need an algorithm that can let us know a bird is flying over Peacetopia as accurately as possible."

- "We want the trained model to take no more than 10sec to classify a new image."
- "We want the model to fit in 10MB of memory."

If you had the three following models, which one would you choose?

Test Accuracy	Runtime	Memory size
97%	1 sec	3MB

Test Accuracy	Runtime	Memory size
99%	13 sec	9MB

Test Accuracy	Runtime	Memory size
97%	3 sec	2MB

Correct

Correct! As soon as the runtime is less than 10 seconds you're good. So, you may simply maximize the test accuracy after you made sure the runtime is <10sec.

3. Based on the city's requests, which of the following would you say is true? **1 / 1 point**

- (•) Accuracy is an optimizing metric; running time and memory size are a satisficing metrics.
- ( ) Accuracy is a satisficing metric; running time and memory size are an optimizing metric.
- ( ) Accuracy, running time and memory size are all optimizing metrics because you want to do well on all three.
- ( ) Accuracy, running time and memory size are all satisficing metrics because you have to do sufficiently well on all three for your system to be acceptable.

✓ **Correct**

1. Has high accuracy
2. Runs quickly and takes only a short time to classify a new image.
3. Can fit in a small amount of memory, so that it can run in a small processor that the city will attach to many different security cameras.

Note: Having three evaluation metrics makes it harder for you to quickly choose between two different algorithms, and will slow down the speed with which your team can iterate. True/False?

- (•) True
- ( ) False

✓ **Correct**

2. After further discussions, the city narrows down its criteria to: **1 / 1 point**

- "We need an algorithm that can let us know a bird is flying over Peacetopia as accurately as possible."
- "We want the trained model to take no more than 10sec to classify a new image."

Train	Dev	Test
3,333,334	3,333,333	3,333,333

(•)

Train	Dev	Test
9,500,000	250,000	250,000

✓ **Correct**

Yes.

5. After setting up your train/dev/test sets, the City Council comes across another 1,000,000 images, called the “citizens’ data”. Apparently the citizens of Peacetopia are so scared of birds that they volunteered to take pictures of the sky and label them, thus contributing these additional 1,000,000 images. These images are different from the distribution of images the City Council had originally given you, but you think it could help your algorithm. **0 / 1 point**

   Notice that adding this additional data to the training set will make the distribution of the training set different from the distributions of the dev and test sets.

   Is the following statement true or false?

   train/dev/test sets. Which of these do you think is the best choice?

- ○

Train	Dev	Test
6,000,000	3,000,000	1,000,000

- ○

Train	Dev	Test
6,000,000	1,000,000	3,000,000

- ○

Train	Dev	Test
3,333,334	3,333,333	3,333,333

- ◉

Train	Dev	Test
9,500,000	250,000	250,000

✓ **Correct**

Yes.

5. After setting up your train/dev/test sets, the City Council comes across another 1,000,000 images, called the “citizens’ data”. Apparently the citizens had originally given you, but you think it could help your algorithm. **0 / 1 point**

Notice that adding this additional data to the training set will make the distribution of the training set different from the distributions of the dev and test sets.

Is the following statement true or false?

"You should not add the citizens' data to the training set, because if the training distribution is different from the dev and test sets, then this will not allow the model to perform well on the test set."

- (•) True
- ( ) False

Incorrect

True is incorrect: Sometimes we'll need to train the model on the data that is available, and its distribution may not be the same as the data that will occur in production. Also, adding training data that differs from the dev set may still help the model improve performance on the dev set. What matters is that the dev and test set have the same distribution.

"You should not add the citizens' data to the training set, because if the training distribution is different from the dev and test sets, then this will not allow the model to perform well on the test set."

- (•) True
- ( ) False

Incorrect

True is incorrect: Sometimes we'll need to train the model on the data that is available, and its distribution may not be the same as

the data that will occur in production. Also, adding training data that differs from the dev set may still help the model improve performance on the dev set. What matters is that the dev and test set have the same distribution.

6. One member of the City Council knows a little about machine learning, and thinks you should add the 1,000,000 citizens' data images to the test set. You object because: **1 / 1 point**

- [x] This would cause the dev and test set distributions to become different. This is a bad idea because you're not aiming where you want to hit.
- [ ] Yes, because having 4.0% training error shows you have high bias.
- [ ] Yes, because this shows your bias is higher than your variance.
- [ ] No, because this shows your variance is higher than your bias.
- [x] No, because there is insufficient information to tell.

✓ **Correct**

8. You ask a few people to label the dataset so as to find out what is human-level performance. You find the following levels of accuracy: **1 / 1 point**

Bird watching expert #1	0.3% error
Bird watching expert #2	0.5% error
Normal person #1 (not a bird watching expert)	1.0% error
Normal person #2 (not a bird watching expert)	1.2% error

If your goal is to have "human-level performance" be a proxy (or estimate) for Bayes error, how would you define "human-level performance"?

- [ ] 0.0% (because it is impossible to do better than this)

Bayes error, how would you define "human-level performance"?

- [ ] 0.0% (because it is impossible to do better than this)
- [x] 0.3% (accuracy of expert #1)
- [ ] 0.4% (average of 0.3 and 0.5)
- [ ] 0.75% (average of all four numbers above)

✓ **Correct**

9. Which of the following statements do you agree with? **1 / 1 point**

- [x] A learning algorithm's performance can be better than human-level performance but it can never be better than Bayes error.
- [ ] A learning algorithm's performance can never be better than human-level performance but it can be better than Bayes error.
- [ ] A learning algorithm's performance can never be better than human-level performance nor better than Bayes error.
- [ ] A learning algorithm's performance can be better than human-level performance and better than Bayes error.
- [ ] 0.0% (because it is impossible to do better than this)
- [x] 0.3% (accuracy of expert #1)
- [ ] 0.4% (average of 0.3 and 0.5)
- [ ] 0.75% (average of all four numbers above)

✓ **Correct**

9. Which of the following statements do you agree with? 1 / 1 point

- (●) A learning algorithm's performance can be better than human-level performance but it can never be better than Bayes error.
- ( ) A learning algorithm's performance can never be better than human-level performance but it can be better than Bayes error.
- ( ) A learning algorithm's performance can never be better than human-level performance nor better than Bayes error.
- ( ) A learning algorithm's performance can be better than human-level performance and better than Bayes error.

Human-level performance	0.1%
Training set error	2.0%
Dev set error	2.1%
Test set error	7.0%

What does this mean? (Check the two best options.)

- [ ] You have underfit to the dev set.
- [ ] You should get a bigger test set.
- [x] You should try to get a bigger dev set.

✓ **Correct**

- [x] You have overfit to the dev set.

✓ **Correct**

seem the most promising to try? (Check two options.)

- [x] Train a bigger model to try to do better on the training set.

Correct

- [ ] Get a bigger training set to reduce variance.
- [ ] Try increasing regularization.
- [x] Try decreasing regularization.

Correct

11. You also evaluate your model on the test set, and find the following: **1 / 1 point**

Human-level performance	0.1%
Training set error	2.0%

the same running time and memory size. However, your system has higher accuracy! However, when Peacetopia tries out your and your competitor's systems, they conclude they actually like your competitor's system better, because even though you have higher overall accuracy, you have more false negatives (failing to raise an alarm when a bird is in the air). What should you do?

- ( ) Look at all the models you've developed during the development process and find the one with the lowest false negative error rate.
- ( ) Ask your team to take into account both accuracy and false negative rate during development.
- (•) Rethink the appropriate metric for this task, and ask your team to tune to the new metric.
- ( ) Pick false negative rate as the new metric, and use this new metric to

drive all further development.

✓ **Correct**

14. You've handily beaten your competitor, and your system is now deployed in Peacetopia and is protecting the citizens from birds! But over the last few

1 / 1 point

✓ **Correct**

- [x] You have overfit to the dev set.

✓ **Correct**

You have only 1,000 images of the new species of bird. The city expects a better system from you within the next 3 months. Which of these should you do first?

- (●) Use the data you have to define a new evaluation metric (using a new dev/test set) taking into account the new species, and use that to drive further progress for your team.
- ( ) Put the 1,000 images into the training set so as to try to do better on these birds.
- ( ) Try data augmentation/data synthesis to get more images of the new type of bird.
- ( ) Add the 1,000 images into your dataset and reshuffle into a new train/dev/test split.

Correct